

Ensemble Pipeline Tutorial

A comprehensive machine learning pipeline for ensemble modeling with ONNX export capabilities and optimization features.

Overview

This pipeline provides a robust framework for training, evaluating, and deploying ensemble machine learning models. It includes features for:

- Model training and evaluation with cost-sensitive optimization
- Hyperparameter tuning using Bayesian optimization
- ONNX model export for deployment
- Interactive configuration through Streamlit UI
- Comprehensive model evaluation and visualization

Installation

1. Clone the repository:

```
git clone https://github.com/QuinnMazaris/ensemble_pipelineV2.git
cd ensemble_pipelineV2
```

2. Set up Python environment:

- **Python Version:** Python 3.12 is recommended for best compatibility
- **Windows Users:** If using Python < 3.12, you must enable long path support
- Create and activate virtual environment:

```
# Windows
python -m venv venv
.\venv\Scripts\activate

# Linux/Mac
python -m venv venv
source venv/bin/activate
```

3. Install dependencies:

```
pip install -r requirements.txt
```

Usage

Running the Pipeline

1. Start the Streamlit interface:

```
streamlit run app/main.py

# Or if looking to host locally

streamlit run app/app.py --server.address 0.0.0.0 --server.port 8501
```

The Streamlit interface will be available at:

- <http://localhost:8501> (local access)
- <http://<your-ip-address>:8501> (network access)

2. Configure settings through the UI or `config.py`:

- Adjust cost weights
- Enable/disable k-fold cross-validation
- Configure feature filtering
- Set optimization parameters
- Choose export format

3. Run the pipeline through the UI or `run.py`.

Model Evaluation

The pipeline provides comprehensive evaluation metrics and visualizations to assess model performance:

1. Performance Metrics:

- **Precision:**
 - Ratio of true positives to all predicted positives ($TP/(TP+FP)$)
 - Measures how many of the predicted defects are actual defects
 - Higher precision means fewer false alarms
- **Recall:**
 - Ratio of true positives to all actual positives ($TP/(TP+FN)$)
 - Measures how many of the actual defects are caught
 - Higher recall means fewer missed defects
- **Accuracy:**
 - Overall correctness $(TP+TN)/(TP+TN+FP+FN)$
 - General measure of model performance
 - Can be misleading with imbalanced data
- **Cost-weighted Performance:**
 - Combines precision and recall using cost weights (C_{FP} and C_{FN})
 - C_{FP} (False Positive Cost): Cost of false alarms (default = 1)
 - C_{FN} (False Negative Cost): Cost of missed defects (default = 30)
 - Total Cost = $(FP \times C_{FP}) + (FN \times C_{FN})$

- Final cost is calculated from the test set's confusion matrix, where FP and FN are the actual counts of false positives and false negatives
- **Confusion Matrix:**
 - Shows detailed breakdown of predictions:
 - True Negatives (TN): Correctly identified good parts
 - False Positives (FP): False alarms (good parts marked as defects)
 - False Negatives (FN): Missed defects
 - True Positives (TP): Correctly identified defects
- **ROC Curve:**
 - Available in Streamlit UI's Model Analysis tab
 - Interactive Plotly visualization
 - Plots True Positive Rate (Recall) vs False Positive Rate
 - Shows trade-off between sensitivity and specificity
 - Area Under Curve (AUC) indicates overall model performance
- **PR Curve** (Precision-Recall):
 - Available in Streamlit UI's Model Analysis tab
 - Interactive Plotly visualization
 - Plots Precision vs Recall
 - Better suited for imbalanced data than ROC
 - Shows trade-off between precision and recall at different thresholds

2. Visualization:

- **Threshold Sweep Plots:**
 - Shows how metrics change across different decision thresholds
 - Identifies optimal thresholds for:
 - Cost minimization (using C_FP and C_FN)
 - Accuracy maximization
 - Helps balance false positives and false negatives
- **Model Comparison Plots:**
 - Side-by-side comparison of all models
 - Shows performance at:
 - Cost-optimal threshold
 - Accuracy-optimal threshold
 - Includes key metrics for each model
- **Class Balance Visualization:**
 - Shows distribution of classes in dataset
 - Displays effect of SMOTE resampling
 - Helps understand class imbalance

Project Structure

```
ensemble_pipelineV2/
├── app/                    # Streamlit application
│   ├── app.py             # Main Streamlit application
│   ├── sidebar.py         # Configuration sidebar
│   └── tabs.py            # Interactive model configuration
```

```

├── utils.py           # Streamlit-specific utilities
├── __init__.py       # Package initialization
├── helpers/          # Core functionality
│   ├── data.py       # Data preparation and processing
│   ├── modeling.py   # Model training and optimization
│   ├── metrics.py    # Evaluation metrics
│   ├── model_export.py # ONNX export functionality
│   ├── plotting.py   # Visualization utilities
│   ├── reporting.py  # Results reporting and analysis
│   ├── utils.py      # General utilities
│   ├── export_metrics_for_streamlit.py # Streamlit metrics export
│   └── __init__.py   # Package initialization
├── data/             # Data directory
├── output/           # Output directory for models, plots, and predictions
├── config.py         # Configuration settings
├── run.py            # Main pipeline execution
├── test_onnx.py      # ONNX export testing
├── requirements.txt  # Project dependencies
└── README.md        # Project documentation

```

Code Architecture

Directory Organization

The project follows a modular architecture with clear separation of concerns:

1. Frontend (**app/** directory):

- Built with Streamlit for interactive web interface
- **app.py**: Main application entry point and tab management
- **sidebar.py**: Configuration UI and settings management
- **tabs.py**: Individual tab implementations for different features
- **utils.py**: Frontend-specific utilities and helper functions
- **__init__.py**: Package initialization

2. Backend (**helpers/** directory):

- Core ML pipeline functionality
- **data.py**: Data loading, preprocessing, and feature engineering
- **modeling.py**: Model training, optimization, and evaluation
- **metrics.py**: Performance metrics and evaluation functions
- **model_export.py**: Model export to ONNX and pickle formats
- **plotting.py**: Visualization and plotting utilities
- **reporting.py**: Results analysis and reporting
- **utils.py**: General utility functions
- **export_metrics_for_streamlit.py**: Metrics export for frontend display
- **__init__.py**: Package initialization and exports

3. Configuration (**config.py**):

- Central configuration file

- Model definitions and hyperparameters
- Pipeline settings and parameters
- File paths and directory structure

4. Data and Output:

- **data/**: Input data directory
- **output/**: Generated files (models, plots, predictions)
 - **models/**: Trained model files
 - **plots/**: Generated visualizations
 - **predictions/**: Model predictions and results - **streamlit_data/**: Model predictions and results for front end

Frontend-Backend Interaction

The application uses a Python-based architecture where Streamlit serves as both frontend and backend:

1. Frontend (Streamlit UI):

- Provides interactive web interface
- Manages user input and configuration
- Displays results and visualizations
- Organized into tabs:
 - Overview: Project summary and status
 - Data Management: Dataset handling
 - Preprocessing Config: Feature engineering settings
 - Model Zoo: Model selection and configuration
 - Model Analysis: Performance metrics and analysis
 - Plots Gallery: Visualization dashboard
 - Downloads: Export model files and results

2. Backend (Python Pipeline):

- Executes ML pipeline operations
- Handles data processing and model training
- Manages model optimization and evaluation
- Exports models and results
- Key components:
 - Data Pipeline: Loading → Preprocessing → Feature Engineering
 - Model Pipeline: Training → Optimization → Evaluation
 - Export Pipeline: Model Conversion → File Export

3. Data Flow:

```
graph TD
    A[User Input (Streamlit UI)] --> B[Configuration (config.py)]
    B --> C[Backend Processing (helpers/)]
```

```
↓  
Results & Models (output/)  
↓  
Display (Streamlit UI)
```

4. Configuration Management:

- Settings stored in `config.py`
- UI updates config through `sidebar.py`
- Backend reads config for pipeline execution
- Changes can be saved permanently to config file

5. Model Management:

- Models defined in `config.py`
- Training handled by `helpers/modeling.py`
- Export managed by `helpers/model_export.py`
- UI configuration through Model Zoo tab

Key Features

1. Modular Design:

- Separate frontend and backend components
- Reusable helper functions
- Clear separation of concerns
- Easy to extend and modify

2. Interactive Configuration:

- Real-time parameter adjustment
- Immediate feedback on changes
- Persistent configuration storage
- Flexible model customization

3. Comprehensive Pipeline:

- End-to-end ML workflow
- Automated data processing
- Model optimization and evaluation
- Results visualization and export

4. Extensible Architecture:

- Easy to add new models
- Customizable preprocessing steps
- Flexible export options
- Configurable evaluation metrics

Configuration

The pipeline is highly configurable through `config.py` and the Streamlit UI. Each setting controls specific aspects of the model training and optimization process:

Model Settings

1. Cost Weights:

- `C_FP` (False Positive Cost): Default = 1
- `C_FN` (False Negative Cost): Default = 30
- These weights are used in the cost-sensitive evaluation of models
- Higher `C_FN` prioritizes reducing false negatives (missed defects)
- Used in threshold optimization to find the optimal decision boundary
- Implementation: `helpers/metrics.py` calculates weighted costs for model evaluation

2. Cross-Validation:

- `USE_KFOLD`: Enable/disable k-fold cross-validation
- `N_SPLITS`: Number of folds (default = 5)
- When enabled:
 - Data is split into `N_SPLITS` stratified folds
 - Models are trained and evaluated on each fold
 - Results are averaged across folds for final metrics
 - Helps assess model stability and generalization
- When disabled:
 - Uses a single 80/20 train-test split
 - Faster training but less robust evaluation
- Implementation: `helpers/modeling.py` handles both split types

3. Feature Filtering:

- `FilterData`: Enable/disable feature filtering
- When enabled, applies two filtering steps:
 1. **Variance Filter** (`VARIANCE_THRESH` = 0.01):
 - Removes features with variance below threshold
 - Eliminates near-constant features
 - Implementation: `helpers/data.py` → `apply_variance_filter()`
 2. **Correlation Filter** (`CORRELATION_THRESH` = 0.95):
 - Removes highly correlated features
 - Keeps one feature from each correlated group
 - Implementation: `helpers/data.py` → `apply_correlation_filter()`
- Both filters are applied sequentially during data preprocessing

Optimization Settings

1. Hyperparameter Optimization:

- `OPTIMIZE_HYPERPARAMS`: Enable/disable Bayesian optimization
- `HYPERPARAM_ITER`: Number of optimization trials (default = 50)
- Uses Optuna for Bayesian optimization:

- Efficiently searches the hyperparameter space
- Optimizes based on cross-validation performance
- Supports early stopping for faster optimization
- Parameter spaces defined in `config.py`:

```
HYPERPARAM_SPACE = {
    'XGBoost': {
        'max_depth': [3,4,5,6],
        'learning_rate': [0.01,0.05,0.1],
        'subsample': [0.6,0.8,1.0],
        'colsample_bytree': [0.6,0.8,1.0],
        'n_estimators': [100,200,400],
        'gamma': [0,0.1,0.2],
    },
    'RandomForest': {
        'n_estimators': [100,200,300],
        'max_depth': [None,5,10],
        'min_samples_split': [2,5,10],
        'min_samples_leaf': [1,2,5],
    }
}
```

- If the user chooses to not use this feature, they are able to adjust the parameters manually via the UI or in `config.py`
- Implementation: `helpers/modeling.py` → `optimize_hyperparams()`

2. Final Model Optimization:

- `OPTIMIZE_FINAL_MODEL`: Enable/disable production model training
- When enabled:
 - If `OPTIMIZE_HYPERPARAMS` is also enabled:
 1. Performs hyperparameter optimization on the full dataset
 2. Uses `HYPERPARAM_ITER` trials (default = 50) to find best parameters
 3. Trains final model with optimized parameters
 - If `OPTIMIZE_HYPERPARAMS` is disabled:
 1. Uses default model parameters from `config.py`
 2. No hyperparameter optimization is performed
 3. Trains final model with default parameters
 - In both cases:
 1. Trains on the entire dataset
 2. Performs cost-sensitive threshold optimization
 3. Saves the production-ready model
- When disabled:
 - Skips final model training completely
 - No production model is created or saved
 - Only cross-validation or single-split models are generated
- Implementation: `helpers/modeling.py` → `FinalModelCreateAndAnalyze()`

3. Class Imbalance Handling:

- **USE_SMOTE**: Enable/disable SMOTE (Synthetic Minority Over-sampling Technique)
- **SMOTE_RATIO**: Ratio of minority to majority class (default = 0.5)
- When enabled:
 - Generates synthetic samples for the minority class
 - Helps prevent model bias towards majority class
 - Applied only to training data, not validation/test
- Implementation: **run.py** applies SMOTE during data preparation
- Note: XGBoost models also use **scale_pos_weight** for additional imbalance handling

Model Export Settings

1. ONNX Export:

- **EXPORT_ONNX**: Enable/disable ONNX model export
- **ONNX_OPSET_VERSION**: ONNX operator set version (default = 12)
- When enabled:
 - Converts models to ONNX format for deployment
 - Preserves feature names and model metadata
 - Supports deployment in production environments
- When disabled:
 - Falls back to pickle format
 - Still maintains all model functionality
- Implementation: **helpers/model_export.py** → **export_model()**

Additional Settings

1. Data Splitting:

- **TEST_SIZE**: Proportion of data for testing (default = 0.2)
- **RANDOM_STATE**: Random seed for reproducibility (default = 42)

2. Output Control:

- **SAVE_MODEL**: Save trained models
- **SAVE_PLOTS**: Generate and save visualizations
- **SAVE_PREDICTIONS**: Save model predictions
- **SUMMARY**: Print detailed progress information

Settings can be modified through the Streamlit UI or directly in **config.py**. Changes made in the UI can be saved permanently to **config.py** using the "Save & Reload Config" button in the sidebar.

Model Export and Deployment

ONNX Export

The pipeline automatically handles model conversion to ONNX format:

1. Export Process:

- Models are converted using `skl2onnx`
- Feature names are preserved for input mapping
- Opset version 12 is used by default
- Automatic fallback to pickle if ONNX unavailable

2. ONNX to Halcon:

- ONNX models can be imported into Halcon
- Input/output tensor names are preserved
- Model metadata includes feature names and thresholds

Export Settings

- Enable ONNX export in the UI or `config.py`
- Adjust opset version if needed (9-15 supported)
- Models are saved in `output/models/`